

```

In [11]: 1 import pandas as pd
          2 import numpy as np
          3 import seaborn as sns
          4 import matplotlib as plt
          5 import warnings
          6 warnings.filterwarnings('ignore')
          7
          8 # For randomized data splitting
          9 from sklearn.model_selection import train_test_split
         10
         11 # To build linear regression_model
         12 import statsmodels.api as sm
         13
         14 # To check model performance
         15 from sklearn.metrics import mean_absolute_error, mean_squared_error
         16
         17 from sklearn.preprocessing import StandardScaler
         18 from sklearn.linear_model import LinearRegression
         19 from statsmodels.compat import lzip
         20 from sklearn import model_selection
         21 from math import sqrt
         22

```

```

In [12]: 1 energy = pd.read_csv("C://Users//E7270//Desktop//Hamoye/energydata_complete.

```

```


In [13]: 1 energy.head()

```

Out[13]:

	date	Appliances	lights	T1	RH_1	T2	RH_2	T3	RH_3	T4	...
0	2016-01-11 17:00:00	60	30	19.89	47.596667	19.2	44.790000	19.79	44.730000	19.000000	...
1	2016-01-11 17:10:00	60	30	19.89	46.693333	19.2	44.722500	19.79	44.790000	19.000000	...
2	2016-01-11 17:20:00	50	30	19.89	46.300000	19.2	44.626667	19.79	44.933333	18.926667	...
3	2016-01-11 17:30:00	50	40	19.89	46.066667	19.2	44.590000	19.79	45.000000	18.890000	...
4	2016-01-11 17:40:00	60	40	19.89	46.333333	19.2	44.530000	19.79	45.000000	18.890000	...

5 rows × 29 columns



In [14]: 1 energy.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 19735 entries, 0 to 19734
Data columns (total 29 columns):
date                19735 non-null object
Appliances          19735 non-null int64
lights              19735 non-null int64
T1                  19735 non-null float64
RH_1                19735 non-null float64
T2                  19735 non-null float64
RH_2                19735 non-null float64
T3                  19735 non-null float64
RH_3                19735 non-null float64
T4                  19735 non-null float64
RH_4                19735 non-null float64
T5                  19735 non-null float64
RH_5                19735 non-null float64
T6                  19735 non-null float64
RH_6                19735 non-null float64
T7                  19735 non-null float64
RH_7                19735 non-null float64
```

In [16]: 1 energy.drop(['date'], axis = 1)

Out[16]:

	Appliances	lights	T1	RH_1	T2	RH_2	T3	RH_3	
0	60	30	19.890000	47.596667	19.200000	44.790000	19.790000	44.730000	19.000000
1	60	30	19.890000	46.693333	19.200000	44.722500	19.790000	44.790000	19.000000
2	50	30	19.890000	46.300000	19.200000	44.626667	19.790000	44.933333	18.920000
3	50	40	19.890000	46.066667	19.200000	44.590000	19.790000	45.000000	18.890000
4	60	40	19.890000	46.333333	19.200000	44.530000	19.790000	45.000000	18.890000
...
19730	100	0	25.566667	46.560000	25.890000	42.025714	27.200000	41.163333	24.700000
19731	90	0	25.500000	46.500000	25.754000	42.080000	27.133333	41.223333	24.700000
19732	270	10	25.500000	46.596667	25.628571	42.768571	27.050000	41.690000	24.700000
19733	420	10	25.500000	46.990000	25.414000	43.036000	26.890000	41.290000	24.700000
19734	430	10	25.500000	46.600000	25.264286	42.971429	26.823333	41.156667	24.700000

19735 rows × 28 columns

In [17]: 1 simple_linear_reg_df = energy[['T2', 'T6']].sample(15, random_state=2)

```
In [18]: 1 energy.isnull().sum()
```

```
Out[18]: date          0
Appliances          0
lights              0
T1                  0
RH_1                0
T2                  0
RH_2                0
T3                  0
RH_3                0
T4                  0
RH_4                0
T5                  0
RH_5                0
T6                  0
RH_6                0
T7                  0
RH_7                0
T8                  0
RH_8                0
T9                  0
RH_9                0
T_out               0
Press_mm_hg         0
RH_out              0
Windspeed           0
Visibility           0
Tdewpoint           0
rv1                  0
rv2                  0
dtype: int64
```

```
In [27]: 1 T2 = x
2 T6 = y
3 rss = np.sum(np.square(y - x))
4 round(rss, 3)
```

```
Out[27]: 3453843.036
```

In [26]: `1 energy.drop(['lights','date'], axis = 1)`


Out[26]:

	Appliances	T1	RH_1	T2	RH_2	T3	RH_3	T4	
0	60	19.890000	47.596667	19.200000	44.790000	19.790000	44.730000	19.000000	4
1	60	19.890000	46.693333	19.200000	44.722500	19.790000	44.790000	19.000000	4
2	50	19.890000	46.300000	19.200000	44.626667	19.790000	44.933333	18.926667	4
3	50	19.890000	46.066667	19.200000	44.590000	19.790000	45.000000	18.890000	4
4	60	19.890000	46.333333	19.200000	44.530000	19.790000	45.000000	18.890000	4
...
19730	100	25.566667	46.560000	25.890000	42.025714	27.200000	41.163333	24.700000	4
19731	90	25.500000	46.500000	25.754000	42.080000	27.133333	41.223333	24.700000	4
19732	270	25.500000	46.596667	25.628571	42.768571	27.050000	41.690000	24.700000	4
19733	420	25.500000	46.990000	25.414000	43.036000	26.890000	41.290000	24.700000	4
19734	430	25.500000	46.600000	25.264286	42.971429	26.823333	41.156667	24.700000	4

19735 rows × 27 columns

In [28]: `1 import seaborn as sns
2 import matplotlib as plt
3 sns.regplot(x="T2", y="T6",
4 data=simple_linear_reg_df)`

Out[28]: `<matplotlib.axes._subplots.AxesSubplot at 0x23444818608>`



In [29]: `1 # independent variables
2 X = energy.drop(["Appliances"], axis=1)
3 # dependent variable
4 y = energy[["Appliances"]]`


```
In [30]: 1 X_train, X_test, y_train, y_test = train_test_split(
2         X, y, test_size=0.30, random_state=42
3     )
```

```
In [32]: 1 X_train.drop(['lights', 'date'], axis = 1)
```

Out[32]:

	T1	RH_1	T2	RH_2	T3	RH_3	T4	RH_4	T5
9129	21.500000	35.626667	17.790000	40.590000	21.700000	35.26	20.39	33.863333	19.600000
2453	19.500000	44.560000	18.700000	44.290000	19.823333	44.50	18.10	43.860000	17.200000
9152	20.790000	35.400000	16.890000	42.030000	21.700000	36.00	19.70	33.200000	19.290000
12694	22.100000	43.260000	19.963333	45.500000	23.390000	39.79	21.10	39.060000	20.660000
16952	24.700000	42.360000	29.856667	31.790000	26.171429	38.59	25.10	39.760000	23.166667
...
11284	21.760000	40.900000	19.390000	43.090000	23.000000	39.00	21.50	39.790000	20.100000
11964	22.390000	43.700000	22.000000	42.066667	22.700000	41.23	21.29	44.120000	19.700000
5390	20.290000	35.700000	18.200000	37.590000	20.100000	37.59	18.20	35.290000	18.500000
860	21.790000	35.560000	20.434000	35.116000	21.200000	36.90	21.10	35.663333	18.050000
15795	21.323333	37.730000	19.890000	38.566667	22.700000	36.59	19.70	39.433333	19.200000

13814 rows × 26 columns



```
In [33]: 1 energy = energy.drop(columns="lights")
```

```
In [45]: 1 #X_train = X_train.drop(columns="date")
2 X_test = X_test.drop(columns="lights")
3 X_test
```

Out[45]:

	T1	RH_1	T2	RH_2	T3	RH_3	T4	RH_4	
8980	20.890000	35.400000	17.760000	39.163333	20.290000	36.900000	19.760000	34.200000	18.
2754	21.890000	53.100000	21.290000	45.360000	21.633333	49.226667	20.533333	40.966667	17.
9132	21.390000	35.500000	17.633333	40.530000	21.666667	35.200000	20.290000	33.760000	19.
14359	21.390000	41.033333	23.890000	34.840000	22.033333	36.933333	22.390000	35.236000	19.
8875	19.963333	35.126667	16.463333	40.126667	20.000000	36.400000	19.260000	34.966667	17.
...
831	22.100000	37.933333	21.196667	35.133333	21.856667	41.100000	21.856667	39.100000	18.
10993	21.700000	38.290000	19.200000	41.560000	23.000000	38.326667	20.260000	39.260000	20.
11761	20.926667	39.400000	18.700000	43.000000	21.790000	37.400000	19.790000	38.290000	19.
12364	22.000000	41.230000	20.700000	41.900000	23.000000	38.790000	21.500000	40.656667	20.
11863	20.890000	42.290000	19.033333	44.963333	21.600000	39.090000	19.700000	41.466667	19.

5921 rows × 26 columns

```
In [39]: 1 print(X_train.head())
```

	T1	RH_1	T2	RH_2	T3	RH_3	T4	RH_4	
9129	21.50	35.626667	17.790000	40.59	21.700000	35.26	20.39	33.863333	
2453	19.50	44.560000	18.700000	44.29	19.823333	44.50	18.10	43.860000	
9152	20.79	35.400000	16.890000	42.03	21.700000	36.00	19.70	33.200000	
12694	22.10	43.260000	19.963333	45.50	23.390000	39.79	21.10	39.060000	
16952	24.70	42.360000	29.856667	31.79	26.171429	38.59	25.10	39.760000	

	T5	RH_5	...	T9	RH_9	T_out	Press_mm_hg	\
9129	19.600000	40.425	...	19.463333	38.260000	0.250000	766.400000	
2453	17.200000	52.000	...	17.200000	46.163333	3.166667	765.266667	
9152	19.290000	39.900	...	19.390000	39.067500	-1.566667	766.000000	
12694	20.660000	58.054	...	20.290000	37.400000	8.833333	753.366667	
16952	23.166667	60.130	...	23.100000	44.466667	21.433333	752.100000	

	RH_out	Windspeed	Visibility	Tdewpoint	rv1	rv2
9129	83.000000	2.000000	65.000000	-2.350000	36.226675	36.226675
2453	85.333333	2.000000	40.000000	0.966667	43.199767	43.199767
9152	89.333333	1.333333	60.666667	-3.100000	24.976055	24.976055
12694	81.000000	1.666667	26.000000	5.733333	16.161125	16.161125
16952	51.000000	2.000000	40.000000	10.800000	17.055346	17.055346

[5 rows x 26 columns]

```
In [46]: 1 linear_model = LinearRegression()
2         #fit the model to the training dataset
3         linear_model.fit(X_train, y_train)
4         #obtain predictions
5         predicted_values = linear_model.predict(X_test)
6         #MAE
7         from sklearn.metrics import mean_absolute_error
8         mae = mean_absolute_error(y_test, predicted_values)
9         round(mae, 3)
```

Out[46]: 53.643

```
In [47]: 1 rss = np.sum(np.square(y_test - predicted_values))
2         round(rss, 3)
```

Out[47]: Appliances 5.191850e+07
dtype: float64

```
In [48]: 1 from sklearn.metrics import mean_squared_error
2         rmse = np.sqrt(mean_squared_error(y_test, predicted_values))
3         round(rmse, 3)
```

Out[48]: 93.64

```
In [49]: 1 from sklearn.metrics import r2_score
2         r2_score = r2_score(y_test, predicted_values)
3         round(r2_score, 3)
```

Out[49]: 0.149

```
In [10]: 1 from sklearn.linear_model import Ridge
2         ridge_reg = Ridge(alpha=0.4)
3         ridge_reg.fit(X_train, y_train)
4
```

```
-----
NameError                                Traceback (most recent call last)
<ipython-input-10-2e0066c0935b> in <module>
      1 from sklearn.linear_model import Ridge
      2 ridge_reg = Ridge(alpha=0.4)
----> 3 ridge_reg.fit(X_train, y_train)

NameError: name 'X_train' is not defined
```

```
In [52]: 1 from sklearn.linear_model import Lasso
2         lasso_reg = Lasso(alpha=0.001)
3
4         lasso_reg.fit(X_train, y_train)
```

Out[52]: Lasso(alpha=0.001, copy_X=True, fit_intercept=True, max_iter=1000,
normalize=False, positive=False, precompute=False, random_state=None,
selection='cyclic', tol=0.0001, warm_start=False)

```
In [ ]: 1 olsmodel_final = sm.OLS(y_train, X_train).fit()
        2 print(olsmodel_final.summary())
```

```
In [ ]: 1 from sklearn.linear_model import Ridge
        2
        3 def get_weights_df(model, feat, col_name):
        4     #this function returns the weight of every feature
        5     weights = pd.Series(model.coef_, feat.columns).sort_values()
        6     weights_df = pd.DataFrame(weights).reset_index()
        7     weights_df.columns = ['Features', col_name]
        8     weights_df[col_name].round(3)
        9     return weights_df
       10
```

```
In [54]: 1 from sklearn.linear_model import Ridge
        2 #ridge_weights_df = get_weights_df(ridge_reg, X_train, 'Ridge_Weight')
        3 lasso_weights_df = get_weights_df(lasso_reg, X_train, 'Lasso_weight')
```

```
-----
NameError                                Traceback (most recent call last)
<ipython-input-54-5df571a1e2d8> in <module>
      1 from sklearn.linear_model import Ridge
      2 #ridge_weights_df = get_weights_df(ridge_reg, X_train, 'Ridge_Weight')
----> 3 lasso_weights_df = get_weights_df(lasso_reg, X_train, 'Lasso_weight')

NameError: name 'get_weights_df' is not defined
```

```
In [55]: 1 #linear_model_weights = get_weights_df(model, x_train, 'Linear_Model_Weight')
        2 #ridge_weights_df = get_weights_df(ridge_reg, x_train, 'Ridge_Weight')
        3 #lasso_weights_df = get_weights_df(lasso_reg, x_train, 'Lasso_weight')
        4
        5 final_weights = pd.merge(linear_model_weights, ridge_weights_df, on='Feature
        6 final_weights = pd.merge(final_weights, lasso_weights_df, on='Features')
```

```
-----
NameError                                Traceback (most recent call last)
<ipython-input-55-c6c0fa20f1ee> in <module>
      3 #lasso_weights_df = get_weights_df(lasso_reg, x_train, 'Lasso_weight')
      4
----> 5 final_weights = pd.merge(linear_model_weights, ridge_weights_df, on='Fe
      6 final_weights = pd.merge(final_weights, lasso_weights_df, on='Features'
      )

NameError: name 'linear_model_weights' is not defined
```

```
In [ ]: 1 from sklearn.metrics import mean_squared_error
        2 rmse = np.sqrt(mean_squared_error(y_test, predicted_values))
        3 round(rmse, 3)
```



```
In [ ]: 1 rss = np.sum(np.square(y_test - predicted_values))
        2 round(rss, 3)
```

```
In [50]: 1 rr = Ridge(alpha=0.01)
        2 rr.fit(X_train, y_train)
        3 pred_train_rr= rr.predict(X_train)
        4 print(np.sqrt(mean_squared_error(y_train,pred_train_rr)))
        5 print(r2_score(y_train, pred_train_rr))
        6
        7 pred_test_rr= rr.predict(X_test)
        8 print(np.sqrt(mean_squared_error(y_test,pred_test_rr)))
        9 print(r2_score(y_test, pred_test_rr))
```

95.21565985217909

```
-----
TypeError                                Traceback (most recent call last)
<ipython-input-50-4518f6ab8d5e> in <module>
      3 pred_train_rr= rr.predict(X_train)
      4 print(np.sqrt(mean_squared_error(y_train,pred_train_rr)))
----> 5 print(r2_score(y_train, pred_train_rr))
      6
      7 pred_test_rr= rr.predict(X_test)
```

TypeError: 'numpy.float64' object is not callable

```
In [53]: 1 model_lasso = Lasso(alpha=0.01)
        2 model_lasso.fit(X_train, y_train)
        3 pred_train_lasso= model_lasso.predict(X_train)
        4 print(np.sqrt(mean_squared_error(y_train,pred_train_lasso)))
        5 print(r2_score(y_train, pred_train_lasso))
        6
        7 pred_test_lasso= model_lasso.predict(X_test)
        8 print(np.sqrt(mean_squared_error(y_test,pred_test_lasso)))
        9 print(r2_score(y_test, pred_test_lasso))
```

95.21569972708544

```
-----
TypeError                                Traceback (most recent call last)
<ipython-input-53-eb1ad2134ad7> in <module>
      3 pred_train_lasso= model_lasso.predict(X_train)
      4 print(np.sqrt(mean_squared_error(y_train,pred_train_lasso)))
----> 5 print(r2_score(y_train, pred_train_lasso))
      6
      7 pred_test_lasso= model_lasso.predict(X_test)
```

TypeError: 'numpy.float64' object is not callable

```
In [ ]: 1
```

